

Lab 4 - Cross-encoder re-ranking

```
In [1]: from helper_utils import load_chroma, word_wrap, project_embedding:  
        from chromadb.utils.embedding_functions import SentenceTransformerE  
        import numpy as np
```

```
In [2]: embedding_function = SentenceTransformerEmbeddingFunction()  
  
        chroma_collection = load_chroma(filename='microsoft_annual_report_',  
        chroma_collection.count())
```

.gitattributes:	1.18k/1.18k [00:00<00:00,
100%	93.2kB/s]
1_Pooling/config.json:	190/190 [00:00<00:00,
100%	21.0kB/s]
README.md:	10.6k/10.6k [00:00<00:00,
100%	1.34MB/s]
config.json:	612/612 [00:00<00:00,
100%	76.0kB/s]
config_sentence_transformers.json:	116/116 [00:00<00:00,
100%	13.6kB/s]
data_config.json:	39.3k/39.3k [00:00<00:00,
100%	4.54MB/s]
pytorch_model.bin:	90.9M/90.9M [00:01<00:00,
100%	73.2MB/s]
sentence_bert_config.json:	53.0/53.0 [00:00<00:00,
100%	5.97kB/s]
special_tokens_map.json:	112/112 [00:00<00:00,
100%	13.9kB/s]
tokenizer.json:	466k/466k [00:00<00:00,
100%	2.60MB/s]
tokenizer_config.json:	350/350 [00:00<00:00,
100%	43.1kB/s]
train_script.py:	13.2k/13.2k [00:00<00:00,
100%	1.72MB/s]
vocab.txt:	232k/232k [00:00<00:00,
100%	1.97MB/s]

modules.json:

349/349 [00:00<00:00,

100%

33.1kB/s]

349

Re-ranking the long tail

Now as we have long tailed documents that are retrieved as relevant documents. We will use Cross Encoder from Sentence Encoder to rank each document according to query:

What is Cross Encoder and How it performs this Ranking thing ? So Basically Sentence Encoder have two types of encoders:

1: Bi-Encoder:

This type of Encoder takes the two Queries separately and pass them from separate encoder and then performs cosine similarity to give the relevancy among both.

2: Cross Encoder:

This type of encoder process inputs (we are considering 2) together as single unit. This allows the model to directly compare and contrast the inputs and understand their relation in better way, also in the end it returns the score from 0 to 1 that shows the similarity among both :)

-Now in short we use Cross Encoder as Ranking thing by giving it Query and Retrieved Documents one by one and in the end we will get score of each document representing how much query is related to that specified document.

```
In [3]: query = "What has been the investment in research and development?"
results = chroma_collection.query(query_texts=query, n_results=10,

retrieved_documents = results['documents'][0]

for document in results['documents'][0]:
    print(word_wrap(document))
    print('')
```

• operating expenses increased \$ 1. 5 billion or 14 % driven by investments in gaming, search and news advertising, and windows marketing. operating expenses research and development (in million \$, except percentages) 2022 2021 percentage change research and development \$ 24, 512 \$ 20, 716 18 % as a percent of revenue 12 % 12 %

Oppt research and development expenses include payroll, employee benefits, stock - based compensation expense, and other headcount - related expenses associated with product development. research and development expenses also include third - party development and programming costs, localization costs incurred to translate software for international markets, and the amortization of purchased software code and services content. research and development expenses increased \$ 3. 8 billion or 18 % driven by investments in cloud engineering, gaming, and linkedin. sales and marketing

```
In [4]: from sentence_transformers import CrossEncoder
cross_encoder = CrossEncoder('cross-encoder/ms-marco-MiniLM-L-6-v2')
```

config.json:	794/794 [00:00<00:00,
100%	91.5kB/s]
pytorch_model.bin:	90.9M/90.9M [00:01<00:00,
100%	73.6MB/s]
tokenizer_config.json:	316/316 [00:00<00:00,
100%	39.0kB/s]
vocab.txt:	232k/232k [00:00<00:00,
100%	3.90MB/s]
special_tokens_map.json:	112/112 [00:00<00:00,
100%	13.9kB/s]

```
In [5]: pairs = [[query, doc] for doc in retrieved_documents]
scores = cross_encoder.predict(pairs)
print("Scores:")
for score in scores:
    print(score)
```

```
Scores:
0.9869341
2.6445777
-0.26802987
-10.731592
-7.706605
-5.6469994
-4.297035
-10.933233
-7.038428
-7.324694
```

```
In [6]: print("New Ordering:")
for o in np.argsort(scores)[::-1]:
    print(o+1)
```

```
New Ordering:
2
1
3
7
6
9
10
5
4
8
```

Re-ranking with Query Expansion

Lastly why not to combine the Query Expansion technique in which we were generating extra queries using LLM that were related to our main query, and Re-ranking to optimize a bit more. But in this scenario we will get 10 documents for each query and then we remove duplicated documents (that were retrieved by every query) and then we can apply re-ranking to get the most relevant documents(may be top 5 or so).

```
In [7]: original_query = "What were the most important factors that contri
generated_queries = [
    "What were the major drivers of revenue growth?",
    "Were there any new product launches that contributed to the in
    "Did any changes in pricing or promotions impact the revenue g
    "What were the key market trends that facilitated the increase
    "Did any acquisitions or partnerships contribute to the revenu
]
```

```
In [8]: queries = [original_query] + generated_queries

results = chroma_collection.query(query_texts=queries, n_results=10)
retrieved_documents = results['documents']
```

```
In [9]: # Deduplicate the retrieved documents
unique_documents = set()
for documents in retrieved_documents:
    for document in documents:
        unique_documents.add(document)

unique_documents = list(unique_documents)
```

```
In [10]: pairs = []
for doc in unique_documents:
    pairs.append([original_query, doc])
```

```
In [11]: scores = cross_encoder.predict(pairs)
```

```
In [12]: print("Scores:")
for score in scores:
    print(score)
```

Scores:

- 9.768024
- 10.0839405
- 5.1418324
- 4.3417664
- 10.042843
- 9.80788
- 8.505109
- 5.27475
- 7.754099
- 9.357721
- 9.918428
- 7.917177
- 7.490655
- 10.000139
- 4.818485
- 1.1369953
- 3.7948635
- 11.0792675
- 4.6518917
- 10.148884
- 10.711212
- 6.9020915
- 3.7681558

```
In [13]: print("New Ordering:")
         for o in np.argsort(scores)[::-1]:
             print(o)
```

New Ordering:

15
22
16
3
18
14
2
7
21
12
8
11
6
9
0
5
10
13
4
1
19
20
17

In []:

In []: